

CLAIM AMOUNT PREDICTIONS FOR JOHNSON INSURANCE PLC

1. Problem Statement

- Building a predictive model that can provide the total amount of claim by a customer in seconds

2. Hypothesis Generation

- a. Policy Deductable:** The monthly contributions on the policy might have the effect on the total amount of the claim by the customer
- b. Incident Hour of The Day:** Some insurance covers regulates the times which are covered by the insurance, this might be one of the factors determining the amount to be claimed
- c. Number of Vehicles involved:** The third party claims will result in higher amounts to be paid to the client, if there is no third party the amount will be less.
- e. Property Damage:** The type of property involved will result in higher claims.
- f. Policy Annual Premium:** The overall contributions by the client might determine how much they client will claim per incident
- g. Umbrella Limit:** Claim limits can determine the number of times that a client can claim with a direct influence on the total claim amount
- h. incident_type :** The type of the incident will determine the severity and the amounts of damages, which will have an impact on whether the total claim is less or high.
- i. collision_type:** Losses will play the same role as capital gains, with direct impact to the amount paid to clients
- j. Months as Customer:** - Loyal customers who have spent many years with the company and contributed more towards their insurance covers could be the ones whose claims are processed quicker
- k. Incident Severity:** The severity of the incident will determine the level of damages which will directly determine the costs and the total claim amount
- j. Authorities Contacted:** Customers who file claims without reporting their incidents to obtain the incident report and case number will determine whether claims will be paid or not.
- k. Witnesses:** Claims that have witnesses can strengthen the filed claim they also guarantee that the claim will be paid and the total amount paid won't have any penalties.
- l. Auto Make:** The make of the vehicle will influence the total amount of the claim to be paid to the client.
- m. Auto Model:** The model of the car will also determine the amount of claim
- n. Auto Year:** The age of the car will affect the amount to be paid to the client, cars that are still 1 to few years old will result in higher claim amounts
- o. Police Report Available:** Claims that do not have the police report might end up not being paid which will result in zero amount being paid, or the amount paid could be higher or accurate without penalties if there is a report

3. Loading Packages and Data

In [2626]:

```
#importing modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings #to ignore warnings
warnings.filterwarnings('ignore');
```

DATA

It consists of the training dataset and testing dataset

Training Data: Will be used to train and test the model accuracy

Testing Data: will be used for submission

In [2627]:

```
#creating dataframes
load_train = pd.read_csv('train.csv', na_values=["?"]) #na_values will replace "?" with NaN
load_test = pd.read_csv('test.csv', na_values=["?"])
load_submission = pd.read_csv('sample_submission.csv')

#making copies to protect the original data
train_original = load_train.copy()
test_original = load_test.copy()
submission_original = load_submission.copy()

train = load_train
test = load_test
sumbission = load_submission
```

4. Understanding the Data

- We check the structure of the train and test data.
- Look at the features present in the datasets


In [2628]:

```
#print first five rows and columns
train.head()
```

Out[2628]:

	Customer_ID	months_as_customer	age	insured_sex	insured_education_level	insured_o
0	Customer_541	239	41	FEMALE	JD	farm
1	Customer_440	108	31	MALE	Masters	prote
2	Customer_482	116	30	MALE	JD	handler
3	Customer_422	8	21	MALE	High School	handler
4	Customer_778	161	38	MALE	PhD	priv-t

5 rows × 37 columns



In [2629]:

```
#shapes of the data
f'train: {train.shape}, test: {test.shape}'
```

Out[2629]:

```
'train: (700, 37), test: (300, 36)'
```

The training data contains 700 observations and the test contains 300 observations

In [2630]:

```
#features of training set
train.columns, len(train.columns)
```

Out[2630]:

```
(Index(['Customer_ID', 'months_as_customer', 'age', 'insured_sex',
        'insured_education_level', 'insured_occupation', 'insured_hobbie
s',
        'insured_relationship', 'capital-gains', 'capital-loss',
        'policy_number', 'policy_bind_date', 'policy_state', 'policy_csl',
        'policy_deductable', 'incident_location', 'incident_hour_of_the_da
y',
        'number_of_vehicles_involved', 'property_damage', 'bodily_injurie
s',
        'policy_annual_premium', 'umbrella_limit', 'insured_zip',
        'incident_date', 'incident_type', 'collision_type', 'incident_seve
rity',
        'authorities_contacted', 'incident_state', 'incident_city', 'witne
sses',
        'police_report_available', 'auto_make', 'auto_model', 'auto_year',
        '_c39', 'total_claim_amount'],
      dtype='object'), 37)
```

In [2631]:

```
#print the data types  
train.dtypes
```

Out[2631]:

Customer_ID	object
months_as_customer	int64
age	int64
insured_sex	object
insured_education_level	object
insured_occupation	object
insured_hobbies	object
insured_relationship	object
capital-gains	int64
capital-loss	int64
policy_number	int64
policy_bind_date	object
policy_state	object
policy_csl	object
policy_deductable	int64
incident_location	object
incident_hour_of_the_day	int64
number_of_vehicles_involved	int64
property_damage	object
bodily_injuries	int64
policy_annual_premium	float64
umbrella_limit	int64
insured_zip	int64
incident_date	object
incident_type	object
collision_type	object
incident_severity	object
authorities_contacted	object
incident_state	object
incident_city	object
witnesses	int64
police_report_available	object
auto_make	object
auto_model	object
auto_year	int64
_c39	float64
total_claim_amount	float64
dtype:	object

In [2632]:

```
train.total_claim_amount
```

Out[2632]:

```
0      14386.67000
1      76440.00000
2      79560.00000
3     121680.00000
4      80640.00000
...
695    106400.00000
696    113733.33000
697     78466.67000
698     97866.67000
699     38400.00000
Name: total_claim_amount, Length: 700, dtype: float64
```

The data consists of the following types

Independent Categorical Data

7-Objects:

- Customer_ID, insured_sex, insured_education_level, insured_occupation, insured_hobbies, insured_relationship
- policy_bind_date, policy_state, incident_location, property_damage, incident_date, incident_type, auto_model, policy_csl
- collision_type, incident_severity, authorities_contacted, incident_state, incident_city, police_report_available, auto_make

Independent Numerical Data

Nominal Data

- number_of_vehicles_involved
- bodily_injuries

13-Int64:

- months_as_customer, age, capital-gains, capital-loss, policy_number, policy_deductable, incident_hour_of_the_day
- policy_annual_premium, umbrella_limit, insured_zip, witnesses, auto_year

3-Float

- _c39
- policy_annual_premium

Dependent/Target Variable

1- float

- total_claim_amount (target variable)

Target Variable

- total_claim_amount (target variable)

5. Exploratory Data Analysis

5.1 Univariables Analysis

We look at the analysis of each variable:


- Categorical: Frequency table, Percentage distribution and Bar-plot
- Numerical: Probability density plot
- **Target Variable: total_claim_amount**

In [2633]:

```
#density distribution of total_claim_amount
plt.figure(1)

plt.subplot(121)
sns.distplot(train.total_claim_amount); #density distribution

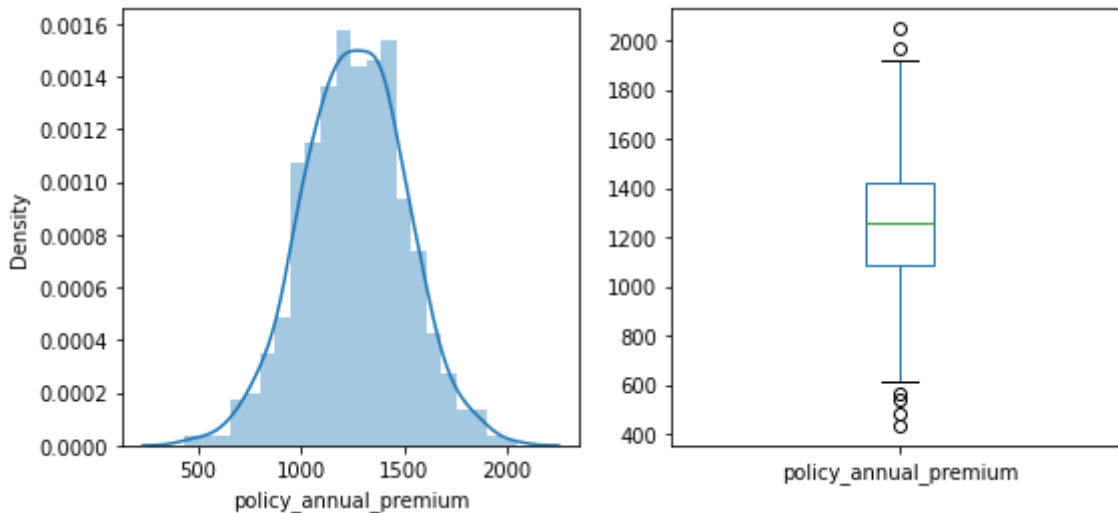
plt.subplot(122)
train.total_claim_amount.plot.box(figsize = (9, 4))
plt.show()
```



- The train.total_claim_amount does not have any extreme values, the distribution is normal, majority of the claims were paid between 50k to 150k.
- **policy_annual_premium**

In [2634]:


```
# distribution of policy_annual_premium
plt.figure(1)
plt.subplot(121)
sns.distplot(train.policy_annual_premium)
plt.subplot(122)
train.policy_annual_premium.plot.box(figsize = (9, 4))
plt.show()
```



- months_as_customer and age

In [2635]:


```
# distribution of months_as_customer
plt.figure(1)
plt.subplot(121)
sns.distplot(train.months_as_customer)
plt.subplot(122)
sns.distplot(train.age)
plt.show()
```



- capital-gains

In [2636]:

```
plt.figure(1)
plt.subplot(121)
sns.distplot(train['capital-gains'])
plt.subplot(122)
train['capital-gains'].plot.box(figsize = (9, 4))
plt.show()
```



In [2637]:

```
train['capital-gains'].head()
```


Out[2637]:

```
0    51400
1         0
2         0
3         0
4    60200
Name: capital-gains, dtype: int64
```

- capital-loss

In [2638]:

```
plt.figure(1)
plt.subplot(121)
sns.distplot(train['capital-loss'])
plt.subplot(122)
train['capital-loss'].plot.box(figsize = (9, 4))
plt.show()
```

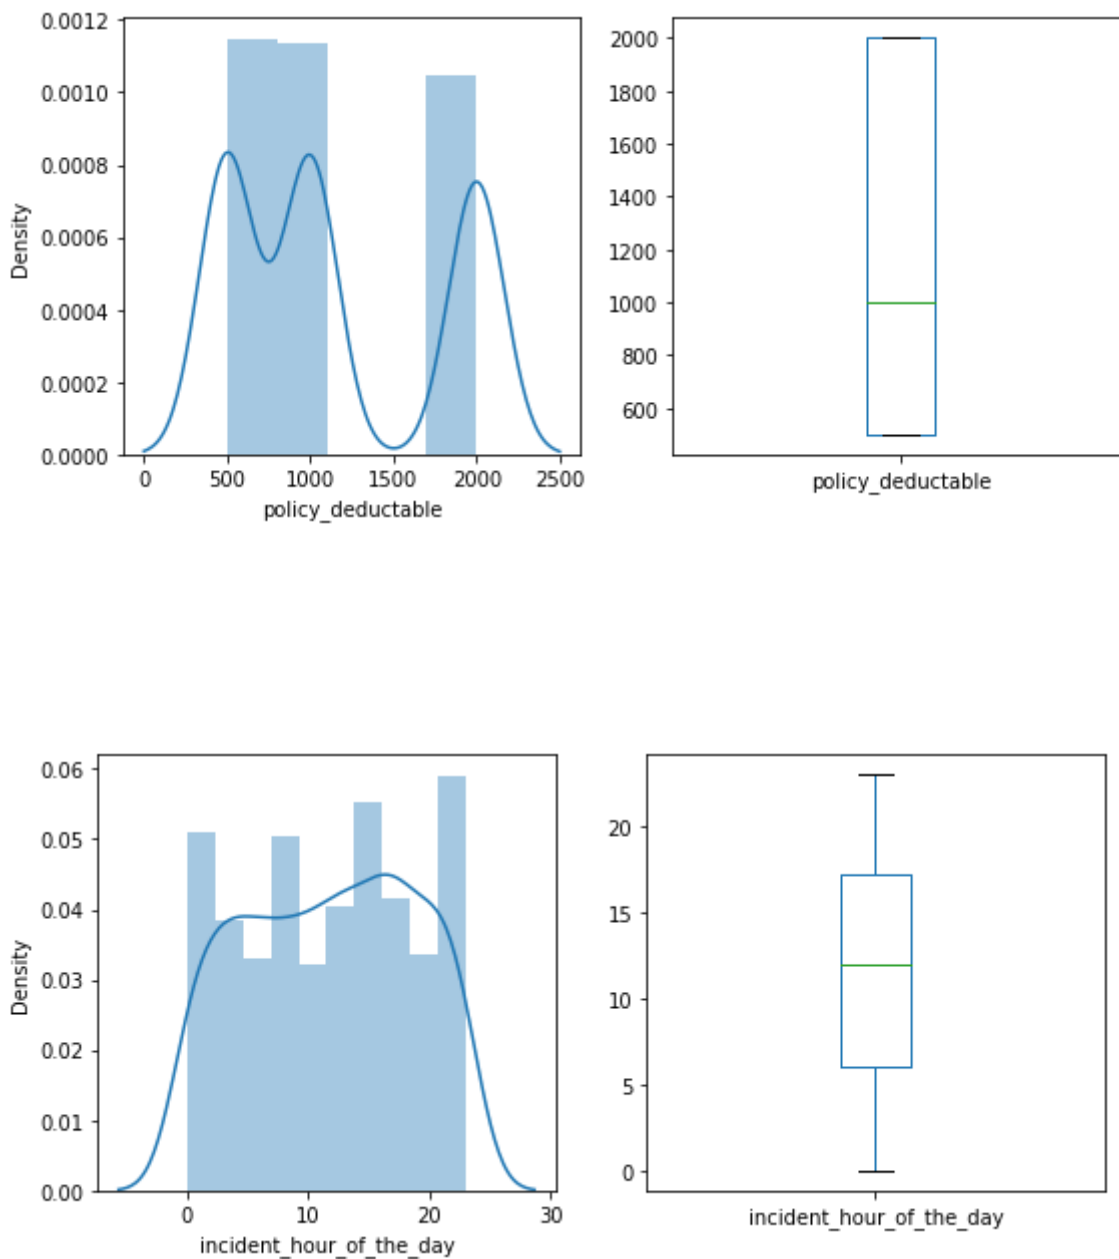


- **policy_deductable** and **incident_hour_of_the_day**

In [2639]:

```
plt.figure(1)
plt.subplot(121)
sns.distplot(train['policy_deductable'])
plt.subplot(122)
train['policy_deductable'].plot.box(figsize = (9, 4))
plt.show()

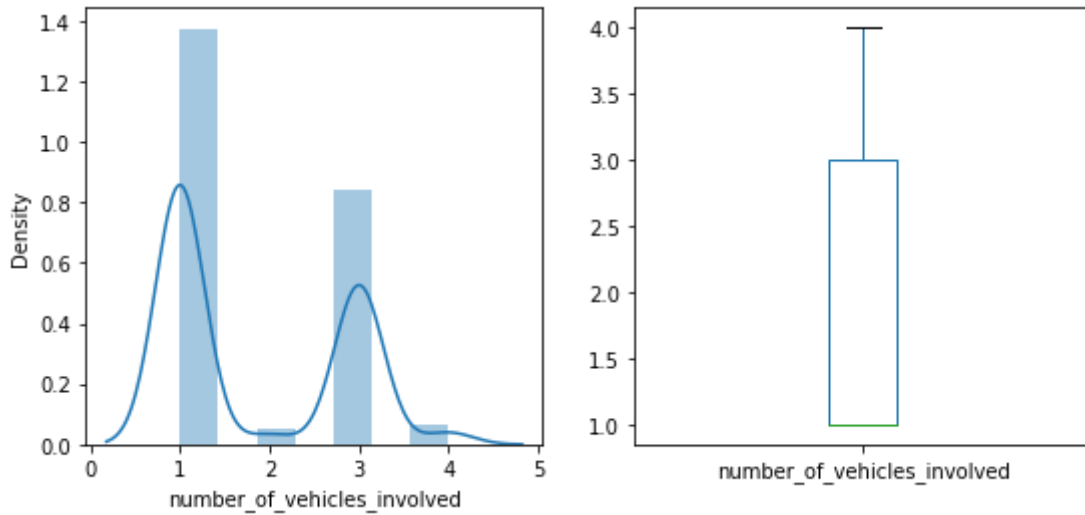
plt.figure(2)
plt.subplot(121)
sns.distplot(train.incident_hour_of_the_day)
plt.subplot(122)
train['incident_hour_of_the_day'].plot.box(figsize = (9, 4))
plt.show()
```



- **number_of_vehicles_involved**

In [2640]:


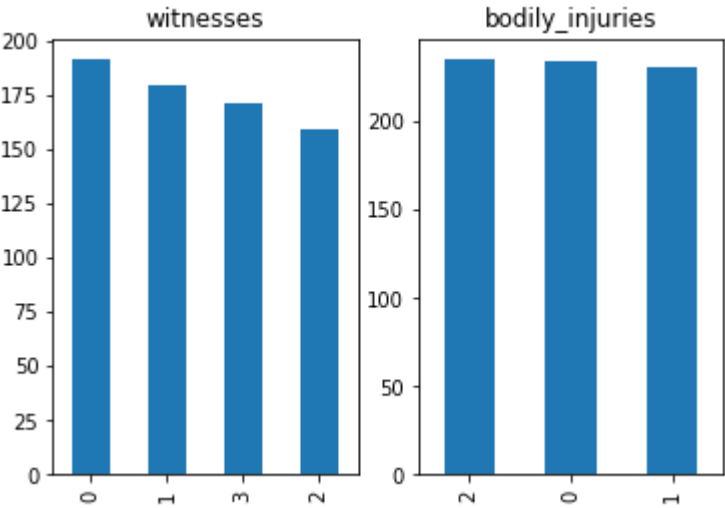
```
plt.figure(2)
plt.subplot(121)
sns.distplot(train.number_of_vehicles_involved)
plt.subplot(122)
train['number_of_vehicles_involved'].plot.box(figsize = (9, 4))
plt.show()
```



- **bodily_injuries**
- **number_of_vehicles_involved, witnesses**
- **authorities_contacted**

In [2641]:

```
#bar plot
plt.figure(1)
plt.subplot(121)
train['witnesses'].value_counts().plot.bar(title = 'witnesses');
plt.subplot(122)
#bar plot
train.bodily_injuries.value_counts().plot.bar(title = 'bodily_injuries')
plt.figure(2)
plt.subplot(121)
train.number_of_vehicles_involved.value_counts().plot.bar(title = 'number_of_vehicles_i
nvolved')
plt.subplot(122)
train.authorities_contacted.value_counts().plot.bar(title = 'authorities_contacted')
plt.show()
```

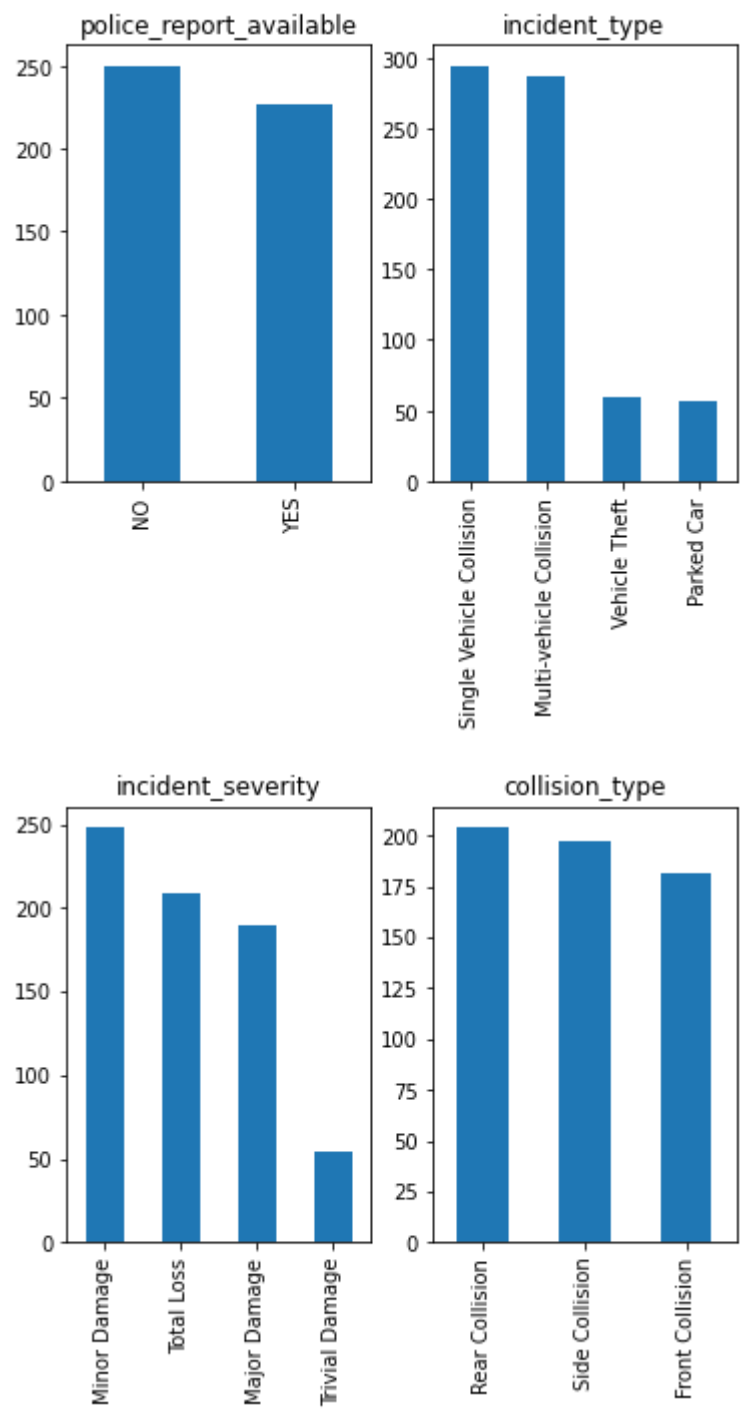


- police_report_available
- police_report_available
- incident_severity
- collision_type

In [2642]:

```
#bar plot
plt.figure(1)
plt.subplot(121)
train['police_report_available'].value_counts().plot.bar(title = 'police_report_availab
le');
plt.subplot(122)
#bar plot
train.incident_type.value_counts().plot.bar(title = 'incident_type')

plt.figure(2)
plt.subplot(121)
train['incident_severity'].value_counts().plot.bar(title = 'incident_severity');
plt.subplot(122)
#bar plot
train.collision_type.value_counts().plot.bar(title = 'collision_type')
plt.show()
```



- auto_make
- insured_sex
- insured_relationship
- insured_education_level

In [2643]:

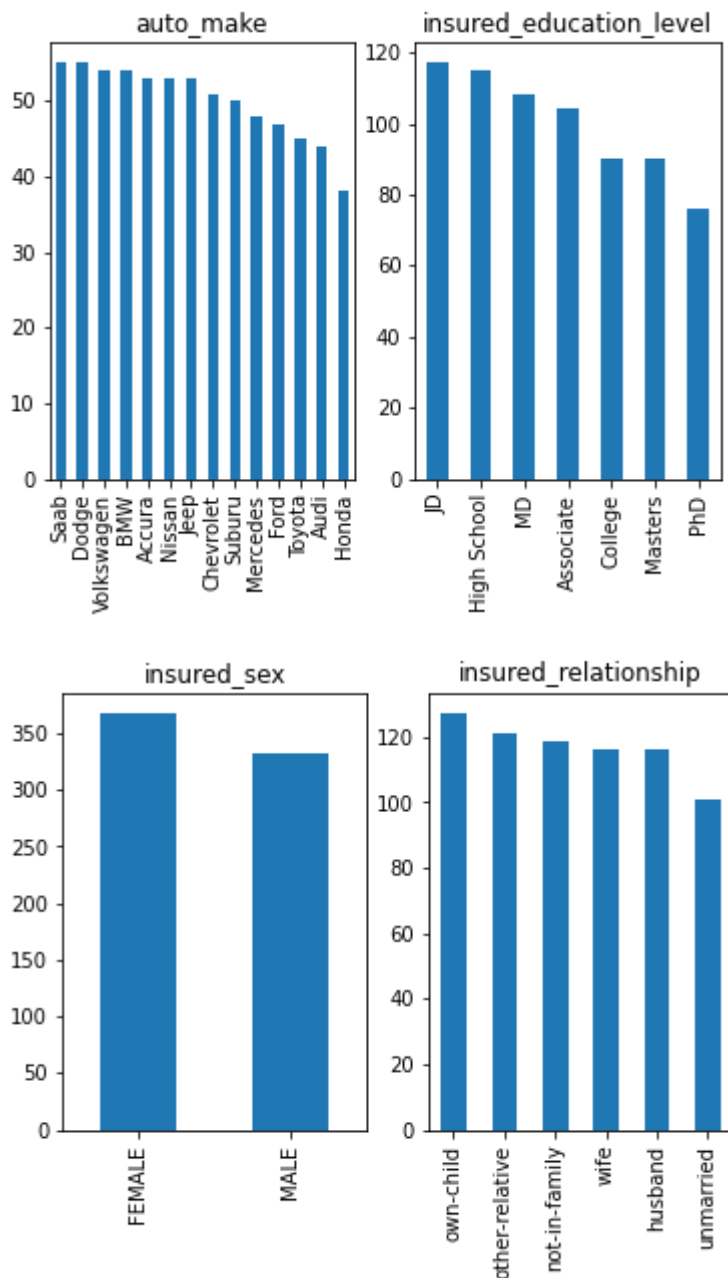
```

#bar plot
plt.figure(1)
plt.subplot(121)
#bar plot
train['auto_make'].value_counts().plot.bar(title = 'auto_make')
plt.subplot(122)
train['insured_education_level'].value_counts().plot.bar(title = 'insured_education_level')
plt.show()

plt.figure(2)
plt.subplot(121)
train.insured_sex.value_counts().plot.bar(title = 'insured_sex')
plt.subplot(122)
train.insured_relationship.value_counts().plot.bar(title = 'insured_relationship')


plt.show()

```



In [2644]:

```
#visualize auto year  
f, ax = plt.subplots(figsize=(10, 5))  
sns.countplot(y="auto_year", data=train, color="c");
```



5.2 Bivariable Analysis


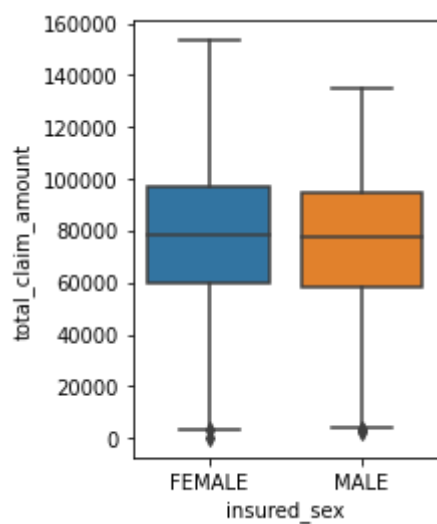
Analysis will be based on Independent variables and target variable to see the relationship between them

- Insured Sex vs Total Claim Amount

In [2645]:

```
plt.figure(1)
plt.subplot(121)
sns.boxplot(x = 'insured_sex', y = 'total_claim_amount', data = train)
plt.show()


plt.figure(2)
plt.subplot(121)
sns.boxplot(x = 'witnesses', y = 'total_claim_amount', data = train);
plt.show()
```



- visualize insured_sex
- total_claim_amount
- insured_education_level

In [2646]:


```
#visualize insured_sex  
#total_claim_amount  
#insured_education_level  
sns.catplot(x="insured_sex", y="total_claim_amount", hue="insured_education_level", kind="bar", data=train);
```



- umbrella_limit
- total_claim_amount

In [2647]:

```
#Using violin to plot the item types as well  
plt.xticks(rotation = 90)  
sns.violinplot(x = train.umbrella_limit, y= train.total_claim_amount);
```




- incident_type
- age
- policy_state

In [2648]:

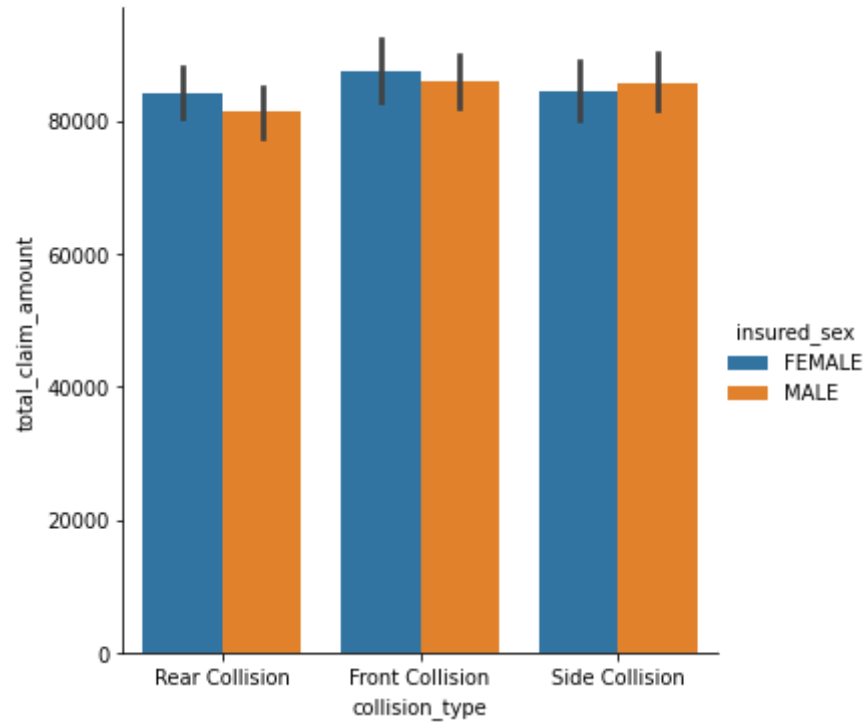

```
#visuallize with Item_Type and Outlet_Size
```

```
sns.catplot(x="policy_state", y="total_claim_amount", hue="insured_sex", kind="box", data=train);
```



In [2649]:

```
#collision type, total_claim_amount, gender  
sns.catplot(x="insured_sex", y="total_claim_amount", hue="collision_type", kind="bar",  
data=train);  
sns.catplot(x="collision_type", y="total_claim_amount", hue="insured_sex", kind="bar",  
data=train);
```



6. Handling Missing Values and Outliers

In [2650]:

```
#sorting missing values columns and filtering 0 missing value cols
missing_value_cols = train.isnull().sum()
missing_value_cols = missing_value_cols[missing_value_cols != 0]
missing_value_cols.sort_values(ascending = False)
```

Out[2650]:

```
_c39                700
property_damage      255
police_report_available  224
collision_type       117
dtype: int64
```

In [2651]:

```
def clean_data(df):
    #Replace missing values with the mode for categorical and or mean/median for numerical
    #train dataset
    df['property_damage'].fillna(df['property_damage'].mode()[0], inplace = True)
    df['police_report_available'].fillna(df['police_report_available'].mode()[0], inplace = True)
    df['collision_type'].fillna(df['collision_type'].mode()[0], inplace = True)

    #drop the _c39 column since it consists of nan only
    df = df.drop(['_c39'], axis = 'columns')

    return df
```

In [2652]:

```
#clean the datasets
train = clean_data(train)
test = clean_data(test)
```

In [2653]:

```
#sorting missing values columns and filtering 0 missing value cols
missing_value_cols = train.isnull().sum()
missing_value_cols = missing_value_cols[missing_value_cols != 0]
missing_value_cols.sort_values(ascending = False)
```

Out[2653]:

```
Series([], dtype: int64)
```

In [2654]:

```
#property damage
train.property_damage.value_counts()
```


Out[2654]:

```
NO      479
YES     221
Name: property_damage, dtype: int64
```

Correlations on Data

In [2655]:

```
#view Correlations
matrix = train.corr()
f, ax = plt.subplots(figsize = (5, 5))
sns.heatmap(matrix, vmax = 0.8, square = True, cmap = 'BuPu');
```




In [2656]:

```
#Compute the correlation matrix
corr = train.corr()
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))
# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(8, 5))
# Generate a custom diverging colormap
cmap = 'nipy_spectral'
# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax= 1, vmin = -1, center=0, square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

Out[2656]:

<AxesSubplot:>



In [2657]:

```
#creates shorth cut for training purposes
train_ = train.copy()
test_ = test.copy()
```

In [2713]:

```
#copy copies for training purposes
train = train_
test = test_
```

In [2714]:

```
#verify if deleted Customer ID is reinstated  
assert 'Customer_ID' in train.columns
```

7. Pre-Processing Data

7.1 One Hot Encoding

In [2715]:

```
def one_hot_encoding(df):  
    #import the library  
    from sklearn.preprocessing import LabelEncoder  
    le = LabelEncoder()  
    #set targets  
    columns_to_fit = [ 'insured_sex', 'insured_occupation', 'insured_relationship',  
                      'policy_state', 'property_damage', 'witnesses', 'police_report_a  
available',  
                      'incident_state', 'auto_model', 'auto_make', 'authorities_contacte  
d', 'months_as_customer',  
                      'insured_hobbies', 'incident_city', 'witnesses', 'insured_educati  
on_level' ]  
    #removed  
    #moved 'collision_type', 'collision_type'  
    #loop df and apply encoding  
    for col in df.columns:  
        if col in columns_to_fit:  
            df[col] = le.fit_transform(df[col])  
    return df
```

In [2716]:

```
train.authorities_contacted.unique()
```

Out[2716]:

```
array([4, 1, 3, 0, 2])
```

In [2717]:

```
#apply one hot encoding  
train = one_hot_encoding(train)  
test = one_hot_encoding(test)
```

In [2718]:

```
f'Education_Level: {train.insured_education_level[0]}'
```

Out[2718]:

```
'Education_Level: 3'
```

7.2 Feature Engineering

In [2719]:

```
def auto_age(df):

    #calculate the years of the auto_year or car
    from datetime import datetime
    current_date_time = datetime.now()
    #replace the YYYY date with number of years
    df['auto_age'] = (current_date_time.year - df['auto_year'])
    #add the years_as_customer
    df['years_as_customer'] = df.months_as_customer//12
    #add the months_out_of_12_as_customer
    df['no_months_as_customer'] = df.months_as_customer%12

    return df
```

In [2720]:

```
#apply feature engineering on the datasets
train = auto_age(train)
test = auto_age(test)
```

In [2721]:

```
#verify processing: years and months of customer 1
f'ID: {train.Customer_ID[0]}, Years: {train.years_as_customer[0]}, Months: {train.no_months_as_customer[0]}'
```

Out[2721]:

```
'ID: Customer_541, Years: 16, Months: 7'
```

In [2722]:

```
def preprocessing_datasets(df):
    #add policy period by the time of the incident
    #use copy df to create it and add to the original df
    dftr = df.copy()
    dftr['incident_date'] = pd.to_datetime(dftr['incident_date'])
    dftr['policy_bind_date'] = pd.to_datetime(dftr['policy_bind_date'])
    df['policy_period'] = (dftr.incident_date - dftr.policy_bind_date)
    #Convert days in policy_period from dtype timedelta to int64
    df['policy_period'] = df['policy_period'].dt.days
    #convert string float fractions in policy_csl into floats
    df_ = df.copy()
    df_[["numerator", "denominator"]] = df_["policy_csl"].str.split("/", expand=True)
    df["policy_csl_float"] = df_["numerator"].astype(float) / df_["denominator"].astype(float)

    return df
```

In [2723]:

```
#convert the dates
train = preprocessing_datasets(train)
test = preprocessing_datasets(test)
```

In [2724]:

```
#confirm processing: policy_period is the number of days of the policy when the accident took place
f'ID: {train.Customer_ID[0]}, Policy period: {test.policy_period[0]} days , Policy CLS: {train.policy_csl_float[0}]'
```

Out[2724]:

'ID: Customer_541, Policy period: 182 days , Policy CLS: 0.5'

7.3 Drop unwanted columns

In [2725]:

```
#apply function to delete columns
def drop_columns(df):

    #set targets
    columns_to_filter = [ 'Customer_ID', 'policy_csl', 'incident_location', 'policy_number', 'policy_bind_date', 'incident_city', 'incident_state', 'incident_date', 'months_as_customer',
                           'auto_year', 'policy_period', 'insured_zip']

    #loop df and delete columns
    for col in df.columns:
        if col in columns_to_filter:
            df = df.drop(col, axis = 1)

    return df
```

In [2726]:

```
#drop unwanted columns
train = drop_columns(train)
test = drop_columns(test)
```

In [2727]:

```
#verify dropped columns
assert 'Customer_ID' not in train.columns
```

In [2728]:

```
#copy data
df_train = train.copy()
df_test = test.copy()
```

8. Building The Model

In [2729]:

```
#drop target variable: total_claim_amount and assign it to y
X = train.drop('total_claim_amount', axis = 1)
y = train.total_claim_amount
```

In [2730]:

```
#Creating Dummies
X = pd.get_dummies(X)
train = pd.get_dummies(train)
test = pd.get_dummies(test)
```

In [2731]:

```
#split the data into train and test
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state =
42)
```

In [2732]:

```
#import the libraries
#import lib and mod
from math import sqrt
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV

#instantiate the model
model = LinearRegression()
#set parameters
parameters = {'fit_intercept':[True,False], 'normalize':[True,False], 'copy_X':[True, F
alse]}
grid = GridSearchCV(model,parameters, cv=None)

#fit/train the model
grid.fit(X_train, y_train);
```

In [2734]:

```
#predict the X_test
y_pred_1 = grid.predict(X_test)
rmse = sqrt(mean_squared_error(y_test, y_pred_1))
print(f'Grid Model RMSE:{rmse}')
```

Grid Model RMSE:20878.17626758436

Model 2

In [2736]:

```
from sklearn.ensemble import GradientBoostingRegressor #For Regression
#Instantiate the model
regress_2 = GradientBoostingRegressor(n_estimators=100, learning_rate=1.0, max_depth=1)
#fit the model
regress_2.fit(X_train, y_train);
```

In [2737]:

```
#prediction with GB
y_pred_2 = regress_2.predict(X_test)

#root mean square error and mean square error of both models
MSE = np.square(np.subtract(y_test,y_pred_2)).mean()
rmse = sqrt(mean_squared_error(y_test, y_pred_2))
print(f'Regress_2 Model RMSE: {rmse}, MSE: {MSE}')
```

Regress_2 Model RMSE: 21789.521358081238, MSE: 474783241.0142782

Model 3: GradientBoostingRegressor and AdaBoostRegressor

In [2738]:

```
#import libraries
from sklearn import ensemble #can also use sklearn.ensemble import GradientBoostRegressor
#set parameters
params = {'n_estimators': 500,
          'max_depth': 4,
          'min_samples_split': 5,
          'learning_rate': 0.01,
          'loss': 'ls'}

#fit GBR
reg_ = ensemble.GradientBoostingRegressor(**params)
reg_.fit(X_train, y_train);
```

In [2739]:

```
#fit adaboost classifier
reg_1=ensemble.AdaBoostRegressor()
reg_1.fit(X_train, y_train);
```

In [2740]:

```
#prediction with GB and AB
reg_pred = reg_.predict(X_test)
reg1_pred = reg_1.predict(X_test)

#root mean square error
rmse = sqrt(mean_squared_error(y_test, reg_pred))
rmse1 = sqrt(mean_squared_error(y_test,reg1_pred))
rmse, rmse1
```

Out[2740]:

(21052.63579053497, 20725.649978770918)

Randomized Search

In [2741]:

```

from sklearn.model_selection import GridSearchCV
from scipy.stats import uniform as sp_rand
from sklearn.linear_model import Ridge
from sklearn.model_selection import RandomizedSearchCV

#instantiate the model
model = LinearRegression()
#set parameters
parameters = {'fit_intercept':[True,False], 'normalize':[True,False], 'copy_X':[True, False]}
grid = GridSearchCV(model,parameters, cv=None)

#fit/train the model
grid.fit(X_train, y_train);

```

In [2744]:

```

# prepare a uniform distribution to sample for the alpha parameter
param_grid = {'alpha': sp_rand()}
# create and fit a ridge regression model, testing random alpha values
model = Ridge()
rsearch = RandomizedSearchCV(estimator=model, param_distributions=param_grid, n_iter=100)
rsearch.fit(X_train, y_train);
print(rsearch)
# summarize the results of the random parameter search
print(rsearch.best_score_)
print(rsearch.best_estimator_.alpha);

```

```

RandomizedSearchCV(cv='warn', error_score='raise-deprecating',
                   estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                                   max_iter=None, normalize=False,
                                   random_state=None, solver='auto',
                                   tol=0.001),
                   iid='warn', n_iter=100, n_jobs=None,
                   param_distributions={'alpha': <scipy.stats._distn_infra
structure.rv_frozen object at 0x261C0170>},
                   pre_dispatch='2*n_jobs', random_state=None, refit=True,
                   return_train_score=False, scoring=None, verbose=0)
0.6618540252818471
0.4678099547970791

```

In [2745]:

```

#prediction with GB and AB
rs_pred = rsearch.predict(X_test)
grid_pred = grid.predict(X_test)
#root mean square error
rmse = sqrt(mean_squared_error(y_test, rs_pred))
rmse1 = sqrt(mean_squared_error(y_test, grid_pred))
rmse, rmse1

```

Out[2745]:

(20774.645553056802, 20878.17626758436)

Polynomial and Linear Regression

In [2746]:

```
#Importing Linear Regression
from sklearn.linear_model import LinearRegression

# Training Model
lm=LinearRegression()
lm.fit(X_train,y_train)
```

Out[2746]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

In [2747]:

```
# importing libraries for polynomial transform
from sklearn.preprocessing import PolynomialFeatures
# for creating pipeline
from sklearn.pipeline import Pipeline

# creating pipeline and fitting it on data
Input=[('polynomial',PolynomialFeatures(degree=2)),('modal',LinearRegression())]

pipe=Pipeline(Input)
pipe.fit(X_train,y_train);
```

In [2748]:

```
#predict wit linear
linear_pred=lm.predict(X_test)
#predict with polynomial
poly_pred=pipe.predict(X_test)
```

In [2749]:

```
print('RMSE for Polynomial Regression=>',np.sqrt(mean_squared_error(y_test,poly_pred)))
print('RMSE for Linear Regression=>',np.sqrt(mean_squared_error(y_test,linear_pred)))
```

RMSE for Polynomial Regression=> 137711.7309840463

RMSE for Linear Regression=> 20756.684446860014

Ordinary Least Squares

In [2750]:

```
#import libraries
import statsmodels.api as sm
# adding a constants
X_train = sm.add_constant(X_train)
X_test = sm.add_constant(X_test)
#train the model
model = sm.OLS(y_train, X_train).fit()
```

In [2751]:

```
#prediction with GB and AB
ols_pred = model.predict(X_test)
#root mean square error
rmse = sqrt(mean_squared_error(y_test, ols_pred))
rmse
```

Out[2751]:

20756.684446091116

Simplified LinearRegression Model

LRM

In [2752]:

```
#import mse
from sklearn.metrics import mean_squared_error

#Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state
=42)
lin_reg_mod = LinearRegression()
#fit the model
lin_reg_mod.fit(X_train, y_train)
```

Out[2752]:

LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

In [2753]:

```
#make a prediction
final_pred = lin_reg_mod.predict(X_test)
#RMSE
rmse = (sqrt(mean_squared_error(y_test, final_pred)))
rmse
```

Out[2753]:

20756.684446860014

In [2754]:

```
final_pred = lin_reg_mod.predict(test)
#Adds total_claim_amout column and fill it
submission = pd.read_csv('sample_submission.csv')
submission['total_claim_amount'] = final_pred
#Adds Customer_ID and fill it
submission['Customer_ID'] = test_original['Customer_ID']
#Convert submission to .csv file format
submission.to_csv('final_pred_model.csv', index = False)
```

Complex Linear Regression Models

In [2755]:

```
#split test and training dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state
=42)

# import necessary libraries
from random import random
from random import randint
from random import seed
from numpy import arange
from numpy import mean
from numpy import std
from numpy import absolute
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RepeatedKFold
from matplotlib import pyplot

# 1. evaluate a model
def evaluate_model(X, y, model):
    # define model evaluation method
    cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
    # evaluate model
    scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error', cv=cv, n_j
obs=-1)
    # force scores to be positive
    return absolute(scores)
```

In [2756]:

```
# 2. plot the dataset and the model's line of best fit
def best_fit(X, y, model):
    #fits the model to the training data
    model.fit(X, y)
    #return the fitted model
    return model

# 3. define the model, evalutate and fit it
model = LinearRegression()
# evaluate model
results = evaluate_model(X_test, y_test, model)
#predict results
print(f'Mean MAE:{mean(results)}, STD: {std(results)}')
# fit the model
fitted_model = best_fit(X_train, y_train, model)
```

Mean MAE:19407.932341490352, STD: 3897.1101134635887

In [2757]:

```
# 3. define the model, evalutate and fit it
model = LinearRegression()
# evaluate model
results = evaluate_model(X_test, y_test, model)
#predict results
print(f'Mean MAE:{mean(results)}, STD: {std(results)}')
# fit the model
fitted_model = best_fit(X_train, y_train, model)
#predict the target
pred_results = fitted_model.predict(X_test)
#RMSE
rmse = (sqrt(mean_squared_error(y_test, pred_results)))
rmse
```

Mean MAE:19407.932341490352, STD: 3897.1101134635887

Out[2757]:

20756.684446860014

Linear Regression with QuantileTransformer

In [2759]:

```
#import libraries
from sklearn.compose import TransformedTargetRegressor
from sklearn.preprocessing import QuantileTransformer

transformer = QuantileTransformer(output_distribution = 'normal')
regressor = LinearRegression()
regression_model_3 = TransformedTargetRegressor(regressor = regressor, transformer = tr
ansformer)
#Model3 with transformed target variable
regression_model_3.fit(X_train, y_train);
```

In [2760]:

```
#predict and measure RMSE
y_pred_3 = regression_model_3.predict(X_test)
rmse = sqrt(mean_squared_error(y_test, y_pred_3))
f'Root Mean Square Error: {rmse}'
```

Out[2760]:

'Root Mean Square Error: 20701.980818671047'

In []: